

Zee Recommender System

```
In [90]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.metrics.pairwise import cosine_similarity
from scipy.sparse import csr_matrix
```

```
In [91]: m = pd.read_fwf('data/zee-movies.dat')
r = pd.read_fwf('data/zee-ratings.dat')
u = pd.read_fwf('data/zee-users.dat')
```

```
In [92]: m.head()
```

```
Out[92]:
```

	Movie ID::Title::Genres	Unnamed: 1	Unnamed: 2
0	1::Toy Story (1995)::Animation Children's Comedy	NaN	NaN
1	2::Jumanji (1995)::Adventure Children's Fantasy	NaN	NaN
2	3::Grumpier Old Men (1995)::Comedy Romance	NaN	NaN
3	4::Waiting to Exhale (1995)::Comedy Drama	NaN	NaN
4	5::Father of the Bride Part II (1995)::Comedy	NaN	NaN

```
In [93]: movies=m.drop(columns=['Unnamed: 1', 'Unnamed: 2'])
```

```
In [94]: movies.head()
```

```
Out[94]:
```

	Movie ID::Title::Genres
0	1::Toy Story (1995)::Animation Children's Comedy
1	2::Jumanji (1995)::Adventure Children's Fantasy
2	3::Grumpier Old Men (1995)::Comedy Romance
3	4::Waiting to Exhale (1995)::Comedy Drama
4	5::Father of the Bride Part II (1995)::Comedy

```
In [95]: movies=movies['Movie ID::Title::Genres'].str.split(':', expand=True)
movies.columns=['MovieID', 'Title', 'Genres']
```

```
In [96]: movies.head()
```

```
Out[96]:
```

	MovieID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

```
In [97]: def findYear(x):  
        if('19' in x):  
            a=x.index('19')  
            return int(x[a:a+4])  
        elif('20' in x):  
            a=x.index('20')  
            return int(x[a:a+4])  
        else:  
            return 0
```

```
In [98]: def Title(x):  
        if('19' in x or '20' in x):  
            return x[:-7]  
        else:  
            return x
```

```
In [99]: movies['Year'] = movies['Title'].apply(lambda x:findYear(x))  
movies['Title'] = movies['Title'].apply(lambda x:Title(x))
```

```
In [100]: movies.head(5)
```

```
Out[100]:
```

	MovieID	Title	Genres	Year
0	1	Toy Story	Animation Children's Comedy	1995
1	2	Jumanji	Adventure Children's Fantasy	1995
2	3	Grumpier Old Men	Comedy Romance	1995
3	4	Waiting to Exhale	Comedy Drama	1995
4	5	Father of the Bride Part II	Comedy	1995

```
In [101]: movies['Title']
```

```
Out[101]: 0          Toy Story
          1          Jumanji
          2    Grumpier Old Men
          3    Waiting to Exhale
          4    Father of the Bride Part II
          ...
          3878    Meet the Parents
          3879    Requiem for a Dream
          3880    Tigerland
          3881    Two Family House
          3882    Contender, The
          Name: Title, Length: 3883, dtype: object
```

```
In [102]: movies.loc[movies['Title'].str.contains('Outlaw, The')]
```

```
Out[102]:
```

	MovieID	Title	Genres	Year
955	967	Outlaw, The	Western	1943

```
In [103]: movies.shape
```

```
Out[103]: (3883, 4)
```

There are 3883 movies in the dataset

```
In [104]: movies['Year'].unique()
```

```
Out[104]: array([1995, 1994, 1996,    0, 1976, 1993, 1992, 1988, 1967, 1964, 1977,
          1965, 1982, 1962, 1990, 1991, 1989, 1937, 1940, 1969, 1981, 1973,
          1970, 1960, 1955, 1956, 1959, 1968, 1980, 1975, 1986, 1948, 1943,
          1950, 1946, 1987, 1997, 1974, 1958, 1949, 1972, 1998, 1933, 1952,
          1951, 1957, 1961, 1954, 1934, 1944, 1963, 1942, 1941, 1953, 1939,
          1947, 1945, 1938, 1935, 1936, 1926, 1932, 1930, 1971, 1979, 1966,
          1978, 1985, 1983, 1984, 1931, 1922,   19, 1927, 1929,   95, 1928,
          1925, 1923, 1999, 1900, 1919, 2000, 1920, 1921], dtype=int64)
```

```
In [105]: mode_year=movies['Year'].mode()
          movies.replace({'Year':{
              0 : mode_year,
              19: mode_year,
              195 : mode_year
          }},
              inplace=True)
```

```
In [106]: movies['Year'].unique()
```

```
Out[106]: array([1995, 1994, 1996, 1976, 1993, 1992, 1988, 1967, 1964, 1977, 1965,
                1982, 1962, 1990, 1991, 1989, 1937, 1940, 1969, 1981, 1973, 1970,
                1960, 1955, 1956, 1959, 1968, 1980, 1975, 1986, 1948, 1943, 1950,
                1946, 1987, 1997, 1974, 1958, 1949, 1972, 1998, 1933, 1952, 1951,
                1957, 1961, 1954, 1934, 1944, 1963, 1942, 1941, 1953, 1939, 1947,
                1945, 1938, 1935, 1936, 1926, 1932, 1930, 1971, 1979, 1966, 1978,
                1985, 1983, 1984, 1931, 1922, 1927, 1929, 1928, 1925, 1923, 1999,
                1900, 1919, 2000, 1920, 1921], dtype=int64)
```

```
In [107]: max(movies['Year']) - min(movies['Year'])
```

```
Out[107]: 100
```

These movies belongs to 19th Century

```
In [108]: r.head()
```

```
Out[108]:
```

	UserID::MovieID::Rating::Timestamp
0	1::1193::5::978300760
1	1::661::3::978302109
2	1::914::3::978301968
3	1::3408::4::978300275
4	1::2355::5::978824291

```
In [109]: ratings=r
ratings.head()
```

```
Out[109]:
```

	UserID::MovieID::Rating::Timestamp
0	1::1193::5::978300760
1	1::661::3::978302109
2	1::914::3::978301968
3	1::3408::4::978300275
4	1::2355::5::978824291

```
In [110]: ratings=ratings['UserID::MovieID::Rating::Timestamp'].str.split(':',expand=True)
ratings.columns=['UserID','MovieID','Rating','Timestamp']
ratings.sample(5)
```

Out[110]:

	UserID	MovieID	Rating	Timestamp
518422	3200	1653	4	968629872
494263	3032	2355	5	970345633
34568	235	1500	5	976770421
324177	1921	455	3	974693607
435841	2664	2145	3	973455388

```
In [111]: ratings.shape
```

Out[111]: (1000209, 4)


The data has total 1000209 ratings

Basic EDA on data

```
In [112]: sns.countplot('Rating',data=ratings)
plt.title('Distribution of movies by Rating')
```

C:\Users\user\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

Out[112]: Text(0.5, 1.0, 'Distribution of movies by Rating')



Most of the times users has rated the movies as '4'

```
In [113]: def decade(x):
            if(x<2000):
                return str((((x)%1900)//10)*10)+'s'
            else:
                return "2000s"
            movies['Decade'] = movies['Year'].apply(lambda x:decade(x))
```

```
In [114]: movies.head()
```

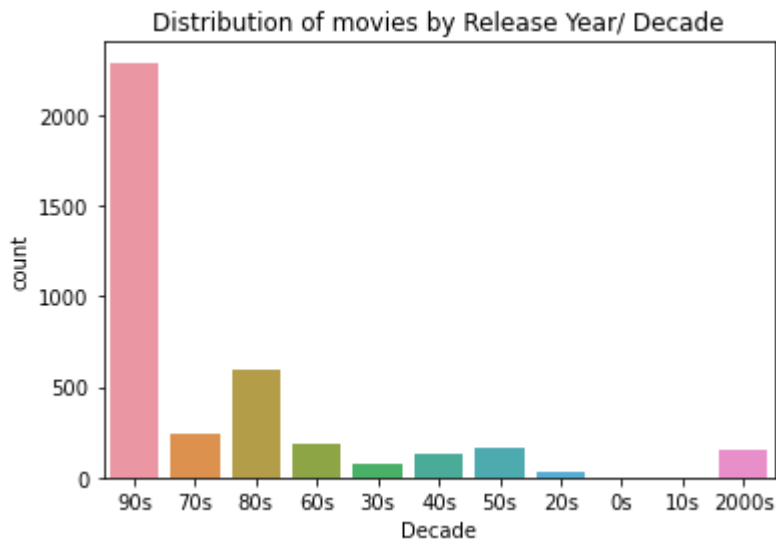
Out[114]:

	MovieID	Title	Genres	Year	Decade
0	1	Toy Story	Animation Children's Comedy	1995	90s
1	2	Jumanji	Adventure Children's Fantasy	1995	90s
2	3	Grumpier Old Men	Comedy Romance	1995	90s
3	4	Waiting to Exhale	Comedy Drama	1995	90s
4	5	Father of the Bride Part II	Comedy	1995	90s

```
In [115]: sns.countplot('Decade',data=movies)
plt.title('Distribution of movies by Release Year/ Decade')
```

C:\Users\user\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

Out[115]: Text(0.5, 1.0, 'Distribution of movies by Release Year/ Decade')



Most of the movies are from 90's

```
In [116]: users=u
u.sample(5)
```

```
Out[116]:
```

	UserID::Gender::Age::Occupation::Zip-code
1014	1015::M::35::3::11220
5682	5683::F::18::9::94538
2926	2927::M::25::12::24060
5144	5145::M::35::7::77565-2332
1270	1271::F::45::6::54401

```
In [117]: users=users['UserID::Gender::Age::Occupation::Zip-code'].str.split(':',expand=True)
users.columns=['UserID', 'Gender', 'Age', 'Occupation', 'ZipCode']
users.sample(5)
```

```
Out[117]:
```

	UserID	Gender	Age	Occupation	ZipCode
1950	1951	F	18	4	90630
3445	3446	M	25	7	30620
423	424	M	25	17	55112
623	624	M	25	1	75207
4219	4220	M	35	0	97225

```
In [118]: users['Age'].unique()
```

```
Out[118]: array(['1', '56', '25', '45', '50', '35', '18'], dtype=object)
```

```
In [119]: users.replace({'Age':{ '1' : 'Under 18', '18': "18-24",
'25': "25-34",
'35': "35-44",
'45': "45-49",
'50': "50-55",
'56': "56+"}}},inplace=True)
```

```
In [120]: users.replace({'Occupation' : {'0': "other",  
    '1': "academic/educator",  
    '2': "artist",  
    '3': "clerical/admin",  
    '4': "college/grad student",  
    '5': "customer service",  
    '6': "doctor/health care",  
    '7': "executive/managerial",  
    '8': "farmer",  
    '9': "homemaker",  
    '10': "K-12 student",  
    '11': "lawyer",  
    '12': "programmer",  
    '13': "retired",  
    '14': "sales/marketing",  
    '15': "scientist",  
    '16': "self-employed",  
    '17': "technician/engineer",  
    '18': "tradesman/craftsman",  
    '19': "unemployed",  
    '20': "writer"}}},inplace=True)
```



```
In [121]: users.sample(10)
```


Out[121]:

	UserID	Gender	Age	Occupation	ZipCode
3068	3069	M	18-24	college/grad student	97470
3616	3617	M	35-44	executive/managerial	49034
5442	5443	M	25-34	programmer	01915
3263	3264	F	25-34	academic/educator	94536
1148	1149	M	25-34	programmer	98103
5877	5878	F	25-34	other	60640
90	91	M	35-44	executive/managerial	07650
1747	1748	M	50-55	academic/educator	04240
5346	5347	M	25-34	doctor/health care	53705
839	840	F	25-34	clerical/admin	02828

```
In [122]: sns.countplot('Gender',data=users)
plt.title('Distribution of users by Gender')
```

C:\Users\user\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

Out[122]: Text(0.5, 1.0, 'Distribution of users by Gender')




most users are male

```
In [123]: sns.countplot('Age',data=users)
plt.title('Distribution of users by Age')
```

C:\Users\user\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

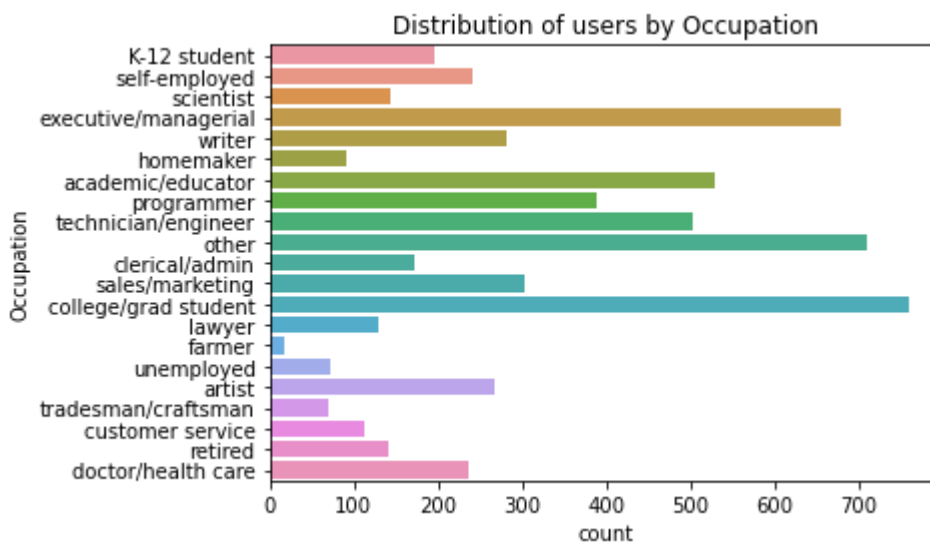
```
Out[123]: Text(0.5, 1.0, 'Distribution of users by Age')
```



Most users belongs to age group of 25-34

```
In [124]: sns.countplot(y='Occupation',data=users,orient="h")
plt.title('Distribution of users by Occupation')
```

```
Out[124]: Text(0.5, 1.0, 'Distribution of users by Occupation')
```



Most number of users are college/ grad student

Merging the datasets

```
In [125]: df1= pd.merge(movies,ratings,how='inner',on='MovieID')
df1.head(10)
```

Out[125]:

	MovieID	Title	Genres	Year	Decade	UserID	Rating	Timestamp
0	1	Toy Story	Animation Children's Comedy	1995	90s	1	5	978824268
1	1	Toy Story	Animation Children's Comedy	1995	90s	6	4	978237008
2	1	Toy Story	Animation Children's Comedy	1995	90s	8	4	978233496
3	1	Toy Story	Animation Children's Comedy	1995	90s	9	5	978225952
4	1	Toy Story	Animation Children's Comedy	1995	90s	10	5	978226474
5	1	Toy Story	Animation Children's Comedy	1995	90s	18	4	978154768
6	1	Toy Story	Animation Children's Comedy	1995	90s	19	5	978555994
7	1	Toy Story	Animation Children's Comedy	1995	90s	21	3	978139347
8	1	Toy Story	Animation Children's Comedy	1995	90s	23	4	978463614
9	1	Toy Story	Animation Children's Comedy	1995	90s	26	3	978130703

```
In [126]: dff = pd.merge(df1,users,how='inner',on='UserID')
dff.sample(20)
```

Out[126]:

	MovieID	Title	Genres	Year	Decade	UserID	Rating	Timestamp
792179	2243	Broadcast News	Comedy Drama Romance	1987	80s	3981	3	9656281
625768	1676	Starship Troopers	Action Adventure Sci-Fi War	1997	90s	5443	3	9599801
71197	2283	Sheltering Sky, The	Drama	1990	90s	795	2	10386931
192639	2901	Phantasm	Horror Sci-Fi	1979	70s	1922	2	9746941
183114	491	Man Without a Face, The	Drama	1993	90s	1861	3	9747051
704093	2757	Frances	Drama	1982	80s	2129	5	9834161
337541	1597	Conspiracy Theory	Action Mystery Romance Thriller	1997	90s	3562	3	9667921
162918	2407	Cocoon	Comedy Sci-Fi	1985	80s	1671	4	10182251
939842	2133	Adventures in Babysitting	Adventure Comedy	1987	80s	4688	4	9636191
345264	555	True Romance	Action Crime Romance	1993	90s	3626	3	9666081
86527	2872	Excalibur	Action Drama Fantasy Romance	1981	80s	963	4	9751201
550776	3219	Pacific Heights	Thriller	1990	90s	5749	2	9628441
610525	590	Dances with Wolves	Adventure Drama Western	1990	90s	4131	3	9653501
218538	1257	Better Off Dead...	Comedy	1985	80s	2109	3	9746561
981853	1230	Annie Hall	Comedy Romance	1977	70s	3241	4	9683461
27166	2136	Nutty Professor, The	Comedy	1963	60s	333	2	9963571
103366	159	Clockers	Drama	1995	90s	1124	3	9749101
851717	527	Schindler's List	Drama War	1993	90s	3069	5	9703581
189800	2193	Willow	Action Adventure Fantasy	1988	80s	1904	3	9747591
772408	1625	Game, The	Mystery Thriller	1997	90s	4783	3	9630051

```
In [127]: dff.tail(50)
```

Out[127]:

	MovieID	Title	Genres	Year	Decade	UserID	Rating
1000159	2028	Saving Private Ryan	Action Drama War	1998	90s	5727	4
1000160	2306	Holy Man	Comedy	1998	90s	5727	2
1000161	2335	Waterboy, The	Comedy	1998	90s	5727	1
1000162	2355	Bug's Life, A	Animation Children's Comedy	1998	90s	5727	4
1000163	2394	Prince of Egypt, The	Animation Musical	1998	90s	5727	3
1000164	2605	Entrapment	Crime Thriller	1999	90s	5727	4
1000165	2634	Mummy, The	Horror	1959	50s	5727	2

```
In [128]: dff.dtypes
```

Out[128]:

MovieID	object
Title	object
Genres	object
Year	int64
Decade	object
UserID	object
Rating	object
Timestamp	object
Gender	object
Age	object
Occupation	object
ZipCode	object
dtype:	object

```
In [129]: dff["Rating"] = dff["Rating"].astype(str).astype(int)
```

```
In [130]: dff['Timestamp'] = pd.to_datetime(dff['Timestamp'],unit='s')
```

```
In [131]: dff.dtypes
```

```
Out[131]: MovieID      object
Title      object
Genres     object
Year       int64
Decade     object
UserID     object
Rating     int32
Timestamp  datetime64[ns]
Gender     object
Age        object
Occupation object
ZipCode    object
dtype: object
```

```
In [132]: dff['Rating'].unique()
```


```
Out[132]: array([5, 4, 3, 2, 1])
```

Even though we have found what is the age group of most number of users, but we are not sure whether the same age group has rated more number of movies. Hence doing this plot on merged Dataset

```
In [192]: sns.countplot('Age', data=dff)
plt.title('Distribution of users by Age')
```

C:\Users\user\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

```
Out[192]: Text(0.5, 1.0, 'Distribution of users by Age')
```

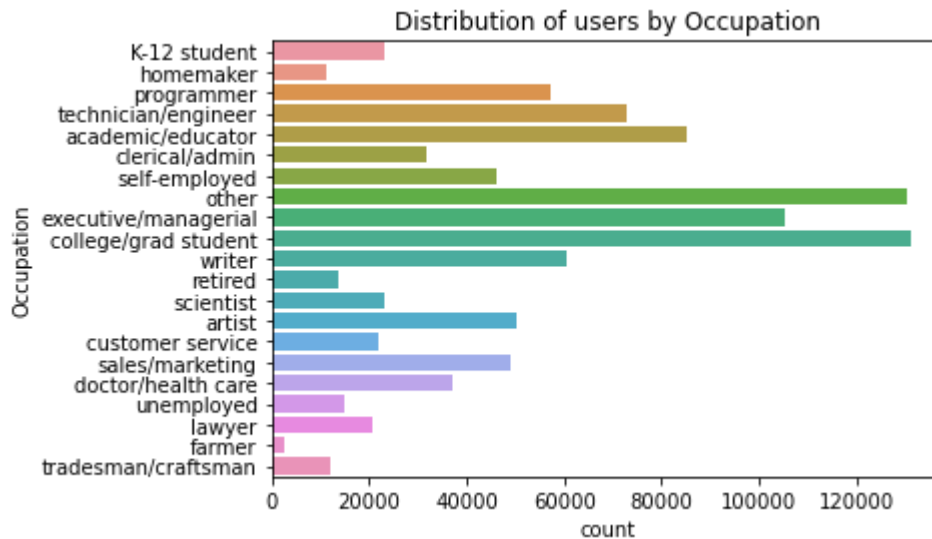


Even though we have found what is the age occupation of most number of users from

users dataset, but we are not sure whether the people of same occupation has rated more number of movies. Hence doing this plot on merged Dataset

```
In [193]: sns.countplot(y='Occupation',data=dff,orient="h")
plt.title('Distribution of users by Occupation')
```

```
Out[193]: Text(0.5, 1.0, 'Distribution of users by Occupation')
```



Grouping the data with respect to Movies to find high rated movies

```
In [133]: temp1=dff.groupby(['MovieID','Title'])\
          .agg(no_of_ratings=('Rating','size'),Avg_Rating=('Rating','mean'))\
          .sort_values(by=['Avg_Rating'],ascending=False)\
          .reset_index()
```

```
In [134]: temp1.loc[temp1['no_of_ratings']>100]
```

Out[134]:

	MovieID	Title	no_of_ratings	Avg_Rating
14	2019	Seven Samurai (The Magnificent Seven) (Shichin...	628	4.560510
15	318	Shawshank Redemption, The	2227	4.554558
16	858	Godfather, The	2223	4.524966
17	745	Close Shave, A	657	4.520548
18	50	Usual Suspects, The	1783	4.517106
...
3623	2555	Baby Geniuses	164	1.701220
3641	2383	Police Academy 6: City Under Siege	149	1.657718
3645	2817	Aces: Iron Eagle III	125	1.640000
3648	3593	Battlefield Earth	342	1.611111
3664	810	Kazaam	120	1.466667

2006 rows × 4 columns

Just considering the cutoff that a movie should be rated by atleast 100 users to be rated as a genuine rated film. Hence **Seven Samurai** is the high rated film in the data

```
In [135]: temp1.loc[temp1['no_of_ratings'] == max(temp1['no_of_ratings'])]
```

Out[135]:

	MovieID	Title	no_of_ratings	Avg_Rating
75	2858	American Beauty	3428	4.317386

More number of users have given their rating for **American Beauty**


```
In [136]: temp1.loc[(temp1['no_of_ratings']>100) & (temp1['Avg_Rating']>=4)]
```

Out[136]:

	MovieID	Title	no_of_ratings	Avg_Rating
14	2019	Seven Samurai (The Magnificent Seven) (Shichin...	628	4.560510
15	318	Shawshank Redemption, The	2227	4.554558
16	858	Godfather, The	2223	4.524966
17	745	Close Shave, A	657	4.520548
18	50	Usual Suspects, The	1783	4.517106
...
373	497	Much Ado About Nothing	667	4.000000
374	1827	Big One, The	102	4.000000
382	2670	Run Silent, Run Deep	220	4.000000
415	1238	Local Hero	351	4.000000
419	2575	Dreamlife of Angels, The (La Vie rêvée des anges)	141	4.000000

302 rows × 4 columns

There are around 302 movies where the movies are considered good with rating above 4

```
In [137]: temp2=dff.groupby(['UserID'])\
            .agg(no_of_ratings=('Rating','size'),Avg_Rating=('Rating','mean'))\
            .sort_values(by=['Avg_Rating'],ascending=False)\
            .reset_index()
```

In [138]: temp2

Out[138]:

	UserID	no_of_ratings	Avg_Rating
0	283	27	4.962963
1	2339	23	4.956522
2	3324	21	4.904762
3	3902	165	4.890909
4	446	51	4.843137
...
6035	5850	58	1.844828
6036	4539	119	1.815126
6037	2744	138	1.304348
6038	4486	51	1.058824
6039	3598	65	1.015385

6040 rows × 3 columns

In [139]: temp2.loc[temp2['no_of_ratings']>10]

Out[139]:

	UserID	no_of_ratings	Avg_Rating
0	283	27	4.962963
1	2339	23	4.956522
2	3324	21	4.904762
3	3902	165	4.890909
4	446	51	4.843137
...
6035	5850	58	1.844828
6036	4539	119	1.815126
6037	2744	138	1.304348
6038	4486	51	1.058824
6039	3598	65	1.015385

6040 rows × 3 columns

The user with ID **283** has given rated movies higher and the user with ID **3598** has rated most of the movies as low

```
In [140]: temp2.loc[temp2['no_of_ratings'] == max(temp2['no_of_ratings'])]
```

Out[140]:

	UserID	no_of_ratings	Avg_Rating
4056	4169	2314	3.551858

The User with UserID **4169** has given ratings for more number of movies

Creating a Pivot Table

```
In [141]: matrix = pd.pivot_table(dff, index="UserID", columns='Title', values='Rating')
matrix.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Index: 6040 entries, 1 to 999
Columns: 3664 entries, $1,000,000 Duck to eXistenZ
dtypes: float64(3664)
memory usage: 168.9+ MB
```

```
In [142]: matrix.head(20)
```

Out[142]:

Title	\$1,000,000 Duck	'Night Mother	'Til There Was You	'burbs, The	...And Justice for All	1-900	10 Things I Hate About You	101 Dalmatians	12 Angry Men	13th Warrior, The	...
UserID											
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
10	NaN	NaN	NaN	4.0	NaN	NaN	NaN	NaN	3.0	4.0	...
100	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	4.0	NaN	NaN	...
1001	NaN	NaN	NaN	NaN	NaN	NaN	NaN	3.0	NaN	NaN	...
1002	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1003	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1004	NaN	NaN	NaN	NaN	NaN	NaN	NaN	4.0	NaN	NaN	...
1005	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1006	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1007	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1008	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1009	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
101	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	4.0	...
1010	NaN	NaN	NaN	1.0	NaN	NaN	3.0	1.0	4.0	NaN	...
1011	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1012	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1013	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1014	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1015	NaN	3.0	2.0	NaN	3.0	NaN	3.0	4.0	NaN	NaN	...

20 rows × 3664 columns




```
In [143]: matrix.fillna(0,inplace=True)
matrix.head(10)
```

Out[143]:

Title	\$1,000,000 Duck	'Night Mother	'Til There Was You	'burbs, The	...And Justice for All	1-900	10 Things I Hate About You	101 Dalmatians	12 Angry Men	13th Warrior, The	...
UserID											
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...
10	0.0	0.0	0.0	4.0	0.0	0.0	0.0	0.0	3.0	4.0	...
100	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...
1000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	4.0	0.0	0.0	...
1001	0.0	0.0	0.0	0.0	0.0	0.0	0.0	3.0	0.0	0.0	...
1002	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...
1003	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...
1004	0.0	0.0	0.0	0.0	0.0	0.0	0.0	4.0	0.0	0.0	...
1005	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...
1006	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...

10 rows × 3664 columns



```
In [144]: matrix.columns
```

Out[144]: Index(['\$1,000,000 Duck', 'Night Mother', 'Til There Was You', 'burbs, The', '...And Justice for All', '1-900', '10 Things I Hate About You', '101 Dalmatians', '12 Angry Men', '13th Warrior, The', ..., 'Young Poisoner's Handbook, The', 'Young Sherlock Holmes', 'Young and Innocent', 'Your Friends and Neighbors', 'Zachariah', 'Zed & Two Noughts, A', 'Zero Effect', 'Zero Kelvin (Kj rlighetens kj tere)', 'Zeus and Roxanne', 'eXistenZ'], dtype='object', name='Title', length=3664)

Item Based Approach : Pearson Correlation

```
In [194]: input_movie = input('Enter a Movie name:')
rating_movie = matrix[input_movie]
similarRated_movies = matrix.corrwith(rating_movie)
```

Enter a Movie name:Liar Liar

```
In [195]: similarRatedMovies
```

```
Out[195]: Title
$1,000,000 Duck          0.063723
'Night Mother           0.048226
'Til There Was You       0.073924
'burbs, The              0.225559
...And Justice for All   0.074494
...
Zed & Two Noughts, A     0.014045
Zero Effect              0.174860
Zero Kelvin (Kj rlighetens kj tere) -0.006114
Zeus and Roxanne         0.067541
eXistenZ                 0.085360
Length: 3664, dtype: float64
```

```
In [196]: df_similar = pd.DataFrame(similarRatedMovies, columns=['Correlation Value'])
df_similar.sort_values('Correlation Value', ascending=False, inplace=True)
df_similar.iloc[1:].head()
```

```
Out[196]:
```

	Correlation Value
Title	
Mrs. Doubtfire	0.499927
Dumb & Dumber	0.459601
Ace Ventura: Pet Detective	0.458654
Home Alone	0.455967
Wedding Singer, The	0.429222

Item Based Approach : Cosine Similarity


Item Similarity Matrix

```
In [197]: item_sim_matrix = cosine_similarity(matrix.T)
item_sim_matrix = pd.DataFrame(item_sim_matrix,index=matrix.columns,columns=matrix.columns)
item_sim_matrix.head()
```

Out[197]:

	Title	\$1,000,000 Duck	'Night Mother	'Til There Was You	'burbs, The	...And Justice for All	1-900	10 Things I Hate About You	101 Dalmatians	101
	Title									
	\$1,000,000 Duck	1.000000	0.072357	0.037011	0.079291	0.060838	0.000000	0.058619	0.189843	0.09
	'Night Mother	0.072357	1.000000	0.115290	0.115545	0.159526	0.000000	0.076798	0.137135	0.1
	'Til There Was You	0.037011	0.115290	1.000000	0.098756	0.066301	0.08025	0.127895	0.128523	0.0
	'burbs, The	0.079291	0.115545	0.098756	1.000000	0.143620	0.000000	0.192191	0.250140	0.1
	...And Justice for All	0.060838	0.159526	0.066301	0.143620	1.000000	0.000000	0.075093	0.178928	0.2

5 rows × 3664 columns




User Similarity Matrix

```
In [198]: user_sim_matrix = cosine_similarity(matrix)
user_sim_matrix = pd.DataFrame(user_sim_matrix,index=matrix.index,columns= matrix.index)
user_sim_matrix.head()
```

Out[198]:

UserID	1	10	100	1000	1001	1002	1003	1004	1005
UserID									
1	1.000000	0.254736	0.123967	0.207800	0.139112	0.110320	0.121384	0.180073	0.103137
10	0.254736	1.000000	0.259052	0.279838	0.158108	0.112659	0.141661	0.431184	0.193049
100	0.123967	0.259052	1.000000	0.306067	0.075625	0.110450	0.358686	0.237292	0.171609
1000	0.207800	0.279838	0.306067	1.000000	0.098971	0.047677	0.201722	0.355619	0.323584
1001	0.139112	0.158108	0.075625	0.098971	1.000000	0.164611	0.053807	0.149848	0.137387

5 rows × 6040 columns



CSR matrix and Printing recommendations based on KNN

```
In [199]: csrmatrix = csr_matrix(matrix.T.values)
csrmatrix
```

Out[199]: <3664x6040 sparse matrix of type '<class 'numpy.float64'>' with 997085 stored elements in Compressed Sparse Row format>

```
In [200]: from sklearn.neighbors import NearestNeighbors
```

```
In [201]: knn = NearestNeighbors(n_neighbors=5,metric='cosine',n_jobs=-1)
knn.fit(csrmatrix)
```

Out[201]: NearestNeighbors(metric='cosine', n_jobs=-1)

```
In [202]: query_movie = input('Enter a movie name:')
distances,indices = knn.kneighbors(matrix[query_movie].values.reshape(1,-1),n_neighbors=5)

for i in range(0,len(distances.flatten())):
    if(i==0):
        print('Recommendations for {0}: \n'.format(query_movie))
    else:
        print('{0}. {1}'.format(i,matrix.columns[indices.flatten()[i]]))
```

Enter a movie name:Liar Liar
Recommendations for Liar Liar:

1. Mrs. Doubtfire
2. Ace Ventura: Pet Detective
3. Dumb & Dumber
4. Home Alone
5. Wayne's World

Matrix Factorization

```
In [154]: #pip install keras
```

```
In [155]: #pip install tensorflow
```

```
In [156]: from keras.models import load_model
from sklearn.model_selection import train_test_split
from keras.layers import Input, Embedding, Flatten, Dot, Dense, Concatenate
from keras.models import Model
```

```
In [174]: u = dff.UserID.unique()
m = dff.MovieID.unique()

userid_to_idx = {o:i for i,o in enumerate(u)}
movieid_to_idx = {o:i for i,o in enumerate(m)}
```

```
In [175]: dff['UserID'] = dff['UserID'].apply(lambda x: userid_to_idx[x])
dff['MovieID'] = dff['MovieID'].apply(lambda x: movieid_to_idx[x])
split = np.random.rand(len(dff)) < 0.8
train_val = dff[split]
test = dff[~split]
split1 = np.random.rand(len(train_val)) < 0.75
train = train_val[split1]
validation = train_val[~split1]
```

```
In [176]: train.shape, validation.shape, test.shape
```

```
Out[176]: ((599763, 12), (200276, 12), (200170, 12))
```

```
In [177]: no_of_unique_movies = len(dff['MovieID'].unique())
no_of_unique_users = len(dff['UserID'].unique())

print(no_of_unique_movies,no_of_unique_users)
```

```
3706 6040
```

```
In [178]: latent_factors = 64
user_input = Input(shape=(1,),name='User-input',dtype='int64')
```

```
In [179]: user_embedding = Embedding(no_of_unique_users,latent_factors,name='user_embedding')
```

```
In [180]: user_vec=Flatten(name="Flatten-Users")(user_embedding)
```

```
In [181]: movie_input = Input(shape=(1,), name="Movie-Input")
movie_embedding = Embedding(no_of_unique_movies, latent_factors, name="Movie-Embedding")
movie_vec = Flatten(name="Flatten-Movies")(movie_embedding)
```

```
In [182]: import tensorflow as tf
import keras
sim =tf.keras.layers.dot([user_vec,movie_vec],name='Similarity-Dot-Product',axes=
model = keras.models.Model([user_input,movie_input],sim)
model.compile(optimizer=tf.keras.optimizers.Adam(lr=1e-4),loss='mse')
```

```
In [183]: model.summary()
```

Model: "model_2"

Layer (type)	Output Shape	Param #	Connected to
User-input (InputLayer)	[(None, 1)]	0	[]
Movie-Input (InputLayer)	[(None, 1)]	0	[]
user_embedding (Embedding)	(None, 1, 64)	386560	['User-input [0][0]']
Movie-Embedding (Embedding)	(None, 1, 64)	237184	['Movie-Inpu t[0][0]']
Flatten-Users (Flatten)	(None, 64)	0	['user_embed ding[0][0]']
Flatten-Movies (Flatten)	(None, 64)	0	['Movie-Embe dding[0][0]']
Similarity-Dot-Product (Dot)	(None, 1)	0	['Flatten-Us ers[0][0]', 'Flatten-Mo vies[0][0]']
Total params: 623,744			
Trainable params: 623,744			
Non-trainable params: 0			

```
In [184]: model_hist = model.fit([train.UserID,train.MovieID],train.Rating,
                                batch_size=128,epochs=20,
                                validation_data = ([validation.UserID,validation.MovieID],
                                verbose=1)
```

```
Epoch 1/20
4686/4686 [=====] - 59s 12ms/step - loss: 14.0455 -
val_loss: 13.9637
Epoch 2/20
4686/4686 [=====] - 61s 13ms/step - loss: 12.4078 -
val_loss: 9.7863
Epoch 3/20
4686/4686 [=====] - 58s 12ms/step - loss: 6.4557 - v
al_loss: 3.9350
Epoch 4/20
4686/4686 [=====] - 58s 12ms/step - loss: 2.7675 - v
al_loss: 2.0785
Epoch 5/20
4686/4686 [=====] - 61s 13ms/step - loss: 1.6608 - v
al_loss: 1.4246
Epoch 6/20
4686/4686 [=====] - 69s 15ms/step - loss: 1.2311 - v
al_loss: 1.1426
Epoch 7/20
4686/4686 [=====] - 65s 14ms/step - loss: 1.0378 - v
al_loss: 1.0070
Epoch 8/20
4686/4686 [=====] - 65s 14ms/step - loss: 0.9410 - v
al_loss: 0.9358
Epoch 9/20
4686/4686 [=====] - 77s 16ms/step - loss: 0.8887 - v
al_loss: 0.8958
Epoch 10/20
4686/4686 [=====] - 74s 16ms/step - loss: 0.8584 - v
al_loss: 0.8721
Epoch 11/20
4686/4686 [=====] - 73s 15ms/step - loss: 0.8394 - v
al_loss: 0.8575
Epoch 12/20
4686/4686 [=====] - 71s 15ms/step - loss: 0.8264 - v
al_loss: 0.8475
Epoch 13/20
4686/4686 [=====] - 73s 16ms/step - loss: 0.8168 - v
al_loss: 0.8397
Epoch 14/20
4686/4686 [=====] - 81s 17ms/step - loss: 0.8088 - v
al_loss: 0.8337
Epoch 15/20
4686/4686 [=====] - 83s 18ms/step - loss: 0.8020 - v
al_loss: 0.8287
Epoch 16/20
4686/4686 [=====] - 70s 15ms/step - loss: 0.7954 - v
al_loss: 0.8241
Epoch 17/20
4686/4686 [=====] - 72s 15ms/step - loss: 0.7894 - v
al_loss: 0.8197
```

```
Epoch 18/20
4686/4686 [=====] - 68s 15ms/step - loss: 0.7834 - v
al_loss: 0.8158
Epoch 19/20
4686/4686 [=====] - 53s 11ms/step - loss: 0.7775 - v
al_loss: 0.8120
Epoch 20/20
4686/4686 [=====] - 53s 11ms/step - loss: 0.7717 - v
al_loss: 0.8079
```

```
In [185]: y_pred = model.predict([test.UserID,test.MovieID],verbose=0)
```

Calculating RMSE

```
In [189]: from sklearn.metrics import mean_squared_error
rmse = mean_squared_error(test.Rating,y_pred,squared=False)
print('Root mean squared error :{:.3f}'.format(rmse))
```


Root mean squared error :0.898

Calculating MAPE

```
In [190]: from sklearn.metrics import mean_absolute_percentage_error
mape = mean_absolute_percentage_error(test.Rating,y_pred)
print('Mean absolute Percentage error :{:.3f}'.format(mape))
```

Mean absolute Percentage error :0.275

```
In [191]: plt.plot(model_hist.history['loss'],'g')
plt.plot(model_hist.history['val_loss'],'b')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train','Test'],loc='upper right')
plt.grid(True)
plt.show()
```



In []: